

Taboo Words in Italian

Francesca Romana Blanda
Tuebingen University

1 Introduction

Language is not merely a tool for neutral information exchange; it is a deeply social and emotional instrument governed by complex cultural norms. Within this landscape, taboo language serves as a primary mechanism for maintaining social order. These prohibited words are not just linguistic anomalies; they reflect the core moral, religious, and ideological systems of a society. In Italy, the "tabooness" of a word is inherently dynamic, shifting across time, location, and the specific "frame" or social context of the interaction.

The persistence of taboo language in human communication can be traced to what researchers call the "Private Information Problem" (PIP). Evolutionarily, humans have developed linguistic "camouflage": structures that can hide vital internal states, such as intense emotions or intentions, to act as a protective barrier in competitive social realities. In the Italian context, this manifests as a unique combination of high emotional arousal and low frequency in formal written language. To navigate the social sanctions associated with these terms, Italian speakers often employ creative linguistic strategies, including phonetic alterations, rhetorical metonymy, and euphemisms.

While global taboos often center on universal themes like sexuality or slurs, their specific offensiveness is deeply bounded to culture. In Italy, for instance, blasphemies carry significant moral weight due to deep seated Catholic roots and the historical relationship between the Church and the State. Furthermore, contemporary Italian discourse frequently targets specific social groups, such as immigrants and nomads, through derogatory language that reinforces identity boundaries.

This paper investigates the semantic patterns and

discursive strategies of offensive language within an Italian context. By utilizing the HaSpeeDe 2 dataset (Sanguinetti et al., 2020) and implementing a language-model-assisted annotation approach, we aim to move beyond simple lexical matching to capture the nuanced, pragmatic nature of "hate spans". Through exploratory data analysis and semantic mapping, we identify recurring attack frames such as dehumanization and demonization, and analyze how they intersect with specific target groups. Ultimately, this study underscores the necessity of a multidimensional, context sensitive approach to understanding the layered nature of offensive language in modern society.

2 Theoretical Framework

In the context of Taboo Language, the word "Taboo" refers to something that is prohibited and limited in its use. In fact, the word taboo is an adaptation of the Polynesian "*tapu*", a term usually translated as forbidden or circumscribed (Graci, 2024). In the Italian cultural framework, Taboo Language is used to give a social type of order to life, in which each individual should be protected from negative or dangerous situations posed by those words (Graci, 2024). The key point is that Taboo words are not just prohibited from the linguistic point of view. Instead, their tabooness gives a guide of conduct which reflects the moral, religious and ideological systems of a certain population (Graci, 2024).

An important concept is also the variability and connotation of taboo words. It is possible to imagine multiple scenarios where the same exact term activates a stronger social and emotional response based on the context compared to another. For example the word "scopare" is used both for cleaning the floor, as well as engaging in sexual intercourse. From an innocent word it can become

a highly taboo word in formal contexts or work environments (Graci, 2024). Goffman introduces the concept of "frame" in order to distinguish the social context that determines whether a word is considered taboo or not (Graci, 2024). Moreover, it should be considered that language is a dynamic variable. Therefore the linguistic prohibition of taboo language is neither universal nor static, as it constantly changes through time, location, and situation (the micro-context) (Graci, 2024).

Analyzing the topic even more in depth, the question of why humans need taboo words in their language arises. Here, (La Porta, 2024) gives a nice answer to this question. The "Tabooness" is born from a biologic and evolutionary necessity of hiding private information. It is called the "Private Information Problem" (PIP): the impossibility of translating mental states that are not conceptual, rather in the realm of emotions, to a propositional type of language. From an evolutionary perspective, the complete transparency of one's own internal states appears to be too dangerous. As a result, humans have developed structures that can hide vital information such as emotions and intentions, through a form of social and linguistic "camouflage". Therefore, Taboo language serves as a kind of protective barrier for the individual to adopt in the reality of a society in a constant competition (La Porta, 2024).

What makes a word "Taboo" in Italy is actually a combination of emotional activation with stable inferences. The psychometric data indicate that Italian taboos are characterized by low valency (unpleasantness), high arousal (physiological activation), and extremely low frequency in the written language (Sulpizio et al., 2024). Linguistically, the strength of a taboo word is not derived from within the word itself, rather from three triggers. The denotation, the defining component (often tied to its etymology), and the connotation. The latter reflects the community's behavior and "encapsulates" past experience contexts, making certain terms unacceptable in formal contexts (Graci, 2024).

In order to reduce the consequences of the utterance of a taboo word in certain social contexts, the creativity of Italian linguistics takes over. In this way, the speaker can distance themselves from the offensive weight of the spoken word (Graci,

2024). Therefore, to reduce the social sanctions associated with a determined taboo word, Italians use omission (signaled by a pause), phonetic alteration (substituting letters or using phonetically similar existent words), or rhetorical techniques such as metonymy and antiphrasis. The same situation occurs when reporting other people's speech: the teller must manipulate the discourse to make sure that they are not identified as the author of the offense (Graci, 2024). Irony can arise from the contrast between the "objective" voice narrating the story and the vulgar language used by the characters, which shows their social inadequacy. In the end, the concept of being "politically correct" acts as a social euphemism, to protect the public image, avoiding being judged as insensitive towards minorities and disabilities (Maher, 2012).

The main categories in Italian are the following: sexual and gender taboo words, blasphemies, illness and death with superstition, and slurs of different types (Graci, 2024). In the first category mentioned, sexual terms are central and often tied to the expressive function of language used to signal shock states or anger. Some examples are "cazzo" (cock) or "figa" (pussy). Additionally, there is a strong gender bias: the most offensive and taboo terms are often derived from female genitalia or female family members. The English translation of these also applies in the same way, with terms such as "figlio di puttana" (motherfucker) and "troia" (cunt) (Graci, 2024). In contrast to many other modern cultures, in Italy, blasphemies belong to a taboo category that has a huge moral weight and is highly sanctioned, because of the deep catholic roots and the close relationship between the Italian Government and the Church. In fact, they constitute a direct attack on the sacred symbols of the Church that reflects the perception of sacrilege as a threat to the moral values of the community (Graci, 2024).

There exist a number of taboos related to unpleasant issues such as incurable diseases (tumor, etc.), often avoided through euphemisms, in response to a "magical thinking", that is afraid of summoning evil just by uttering it. Finally, the last predominant category is the one of the so called "slurs". They are words created against specific social groups and minorities that differ from traditional norms. In this way, they make those

differences even more marked by pointing out what they believe does not belong to their own culture (Sulpizio et al., 2024). In Italy, at the same time, another phenomenon occurs in order to avoid being seen as unsensitive. Certain terms used to denote disabilities or humble jobs are replaced by longer descriptions which are usually less direct in expression. Some examples are "collaboratore scolastico" (lit. scholastic collaborator) instead of "bidello" (janitor) and "non vedente" (person that cannot see) instead of "cieco" (blind person) (Graci, 2024).

Overall it is necessary to note that even though there are common grounds for taboo words in different languages such as sex, or slurs, the offensiveness and perception of those words are specific to each culture. For example, words that are taboo in one language, for example "sex" or "vagina", may be significantly less offensive in another, while still being socially restricted in certain contexts. In the end, the offensiveness of a word in Italian does not depend a lot from its frequency, but more from the sociocultural sensitivity specific to the Italian community (Sulpizio et al., 2024).

3 Methodology

3.1 Data Research and Integration

The dataset utilised by this study has not been obtained ex novo, since it has been derived from an already existing resource from extensive previous work, which is hereby cited (Sanguinetti et al., 2020). In more detail, the dataset that was used belonged to a campaign of evaluation called "HaSpeede 2" which included both *news* and *tweets* data. They are both divided into two main tasks:

- Offensiveness classification and presence of stereotypes
- identification of the *hate spans*

The datasets were downloaded from a public repository in tabular format: *TSV* for the classification tasks and *TXT* for the sequence labeling task. The different components have been integrated successfully in one unique dataset afterwards. Therefore the final dataset combines all of the following:

1. The classification task datasets (Task A/B) contain information about offensiveness and stereotype on a sentence level, for both *tweets* and *news*.

2. The *hate spans* datasets (Task C) provide token-level annotations in BIO format, for both *tweets* and *news* as well.

In order to obtain a unified representation, a *merging* process was executed between the tasks, associating to each text the offensive part (offensive span), reconstructed starting from the BIO labels. The token sequences labeled as offensive: B and I, were merged to form a continuous string, meaning the offensive span. Successively, the *news* and *tweets* datasets were concatenated in one single corpus adding also one column specifying the language (Italian).

3.2 Preprocessing and Annotation

The original labels were renamed in order to guarantee coherence and compatibility with the datasets for Korean and Sinhala analyzed by this project. To be exact, the changes made were the following:

1. The variable *offensiveness* was switched to binary with the values "OFF" for an offensive entry and "NOT" for a not offensive one
2. The variable *stereotype* was transformed into *target_type*, with the distinction between:
 - "GRP": group target
 - "IND": individual target
 - "null": absence of offensiveness and stereotype

Moreover, the missing values were properly handled as well as the labeling convention, to ensure a smooth comparison between the datasets later on.

3.3 Annotation of the semantic categories

In order to enrich the dataset with semantic information, the variable *target_group_attribute* was introduced, which identifies the social category or identity that is the target of the offensive expression. Some examples of this are: ethnicity, religion, gender, sexual orientation, etc. Given the complex and contextual nature of this type of classification, a language model assisted annotation approach was adopted. In particular:

1. A subset of the dataset containing the text (context) and an empty column for annotation was exported.
2. This file was provided to a model based annotator (ChatGPT), which was instructed to

classify each example based on the overall meaning of the sentence, avoiding simple lexical matching.

3. The annotated result was later reimported and integrated into the main dataframe.

In parallel, the *target_group* variable has also been added, which more concretely specifies the target group (e.g. “immigrants”, “women”, “LGBTQ+”). Some manual corrections and normalizations were also made, such as the standardization of labels (e.g. “lgbtq+” to “LGBTQ+”); correcting specific cases through pattern matching (for example occurrences of “rom” reclassified as “Nomads”).

3.4 Exploratory Data Analysis (EDA)

An exploratory data analysis phase was conducted to understand the distribution of the annotated variables and verify the internal consistency of the enriched dataset. The analysis was implemented using standard data analysis libraries in Python, specifically *pandas* for data manipulation and *matplotlib* and *seaborn* for visualization.

3.4.1 Distribution of *target_group_attribute*

To analyze the distribution of semantic categories (*target_group_attribute*), the absolute frequency of each category was calculated through aggregation operations on the dataframe. In particular, a count function (e.g. *value_counts()* in *pandas*) was used and applied to the corresponding column.

The resulting frequencies were then sorted in descending order to highlight the most represented categories in the dataset. The results were displayed using bar charts (bar plots), where:

- The x-axis represents the *target_group_attribute* categories;
- The y-axis represents the frequency (number of occurrences);
- The categories were eventually rotated on the horizontal axis to improve its readability.

3.4.2 Cross-analysis between *target_group* and *target_group_attribute*

To explore the relationship between the specific target group (*target_group*) and the general semantic category (*target_group_attribute*), a contingency analysis was performed.

In particular, a double entry frequency table was constructed using a cross tabulation function (e.g.

pd.crosstab()), which calculates the number of occurrences for each combination of the two variables. The resulting table was used to identify dominant associations between specific groups and semantic categories and check the consistency of the annotations (for example, that specific groups such as “women” are predominantly associated with the category “gender”). For visualization, two possible strategies have been adopted:

1. Grouped bar charts (grouped bar plots), where the distributions of semantic categories are shown for each *target_group*;
2. Heat maps (heatmaps), where the intensity of the colour represents the frequency of the combination.

3.4.3 Consistency display and verification

The visualizations produced were used not only for descriptive purposes, but also as a qualitative validation tool for the annotations. In fact, highly unbalanced or unexpected distributions were analyzed manually. Moreover, any inconsistencies between *target_group* and *target_group_attribute* were identified by inspecting the less frequent cells of the contingency table and some ambiguous cases have been reviewed to verify the correctness of the labeling. This process improved the overall quality of the dataset, reducing systematic errors and increasing the semantic consistency of the annotations.

3.5 Alternative method (Appendice)

In a preliminary phase of the research, an alternative approach for semantic category annotation was considered, based on more automated methods (e.g. rule based classification or traditional NLP models). However, this approach was later discarded. The main critical issues encountered included:

- Difficulty in capturing the contextual meaning of offensive expressions;
- Low accuracy in identifying semantic categories, especially in the presence of implicit or metaphorical language;
- Strong dependence on specific lexical choices, which cannot be generalized to all linguistic variants of offensive language.

Given the relatively small size of the dataset, it was considered more effective to adopt a language

model assisted annotation approach, capable of interpreting semantic content more flexibly and contextually. This led to qualitatively better results than the alternative method initially considered.

4 Results: Quantitative and Qualitative Analysis

The results of the Exploratory Data Analysis (EDA) on the enriched dataset reveal specific trends in the Italian taboo landscape, moving from broad distribution patterns to specific linguistic and semantic features.

4.1 Global Distribution and Lexical Density

The first stage of analysis confirms the structural composition of the dataset. We observe a significant class imbalance that characterizes the detection of offensive language as a "sparse signal" task.




Figure 1: Comparison of sentence lengths between of-fensive (OFF) and non-offensive (NOT) classes.

Normalized class distribution:	
target_group_attribute	
Immigrants	0.659104
Nomads	0.151446
Religion	0.121384
Gender & Sexual Orientation	0.041974
Miscellaneous	0.014748
Political Affiliation	0.009075
Socioeconomic status	0.001702
Race, Ethnicity & Nationality	0.000567

Figure 2: Percentage breakdown of the various target groups within the corpus.

As seen in Figure 1, there is no drastic difference in length between offensive and neutral sentences, suggesting that offensiveness is not a byproduct of sentence brevity. Figure 2 further quantifies the dominance of specific groups, with "Immigrants" and "Nomads" representing the vast majority of the normalized distribution.

4.2 Semantic Mapping and Similarity

To understand how these taboo terms relate to one another in a vector space, we analyzed the semantic similarity of the hate spans.




Figure 3: Cluster analysis of target groups (e.g., Islam, Christianity, LGBTQ+) based on linguistic similarity.




Figure 4: Higher-level semantic clustering of group attributes like Religion and Race.

Figure 3 shows that "Immigrants," "Nomads," and "Islam" form a central semantic cluster, suggesting that in Italian discourse, these identities are often targeted using overlapping linguistic frames. Figure 4 reinforces this by showing that "Race" and "Religion" are the most semantically dense attributes in the corpus.

4.3 Category Prevalence and Morphological Complexity

We examined the frequency of each category and the linguistic effort (sentence length) associated with them.




Figure 5: Raw count of instances per category, highlighting the prevalence of Immigrants and Religion.

Mean sentence length per category:	
target_group_attribute	
Gender & Sexual Orientation	32.716216
Immigrants	26.733219
Miscellaneous	26.153846
Nomads	31.599251
Political Affiliation	36.562500
Race, Ethnicity & Nationality	14.000000
Religion	32.500000
Socioeconomic status	31.000000

Figure 6: Average word count per sentence across different taboo categories.

The count in Figure 5 highlights the heavy focus on "Immigrants" (over 1,100 instances). Interestingly, Figure 6 shows that "Political Affiliation" and "Gender" categories result in longer mean sentence lengths (over 30 words), indicating that these offenses are often embedded in more complex argumentative structures than religious or racial slurs.

4.4 Intersectionality and Attack Frames

Finally, we analyzed how specific identities intersect with various types of linguistic attacks.




Figure 7: Correlation matrix between specific target groups and their associated attributes.




Figure 8: Mapping of target attributes against the style of attack (e.g., Insult, Dehumanization).

The final correlation matrices provide the most granular view of the results. Figure 7 shows an absolute correlation (1162 instances) between the "Immigrant" group and its corresponding attribute. More crucially, Figure 8 illustrates that "Insult" is the most universal attack frame across all attributes, while "Dehumanization" and "Demonization" are more selectively applied to "Immigrants" and "Religion," respectively, confirming our hypothesis on the cultural specificities of Italian taboos.

The data analysis allowed us to obtain a detailed view of the distribution and semantic characteristics of offensive language within the Italian dataset considered. For example, examining the *target_group_attribute* variable highlights a non-uniform distribution of semantic categories. Some frames are significantly more represented than others, suggesting the presence of recurring patterns in the construction of linguistic offense. In particular, categories related to social identity such as ethnicity, gender, and religion, emerge more frequently, along with discursive strategies such as dehumanization and pathologization. This finding is in line with the literature on taboo language, according to which such expressions perform functions not only emotional, but also social and identity related. Conversely, less frequent categories tend to reflect more specific or highly context dependent forms of attack, suggesting lower productivity or greater pragmatic variability.

A further level of analysis concerns the relationship between *target_group* and *target_group_attribute*. Cross-analysis between these two variables highlighted a strong correlation between specific target groups and certain semantic frames. Defined social groups, such as ethnic

or religious ones, are most frequently associated with frames of demonization or representations of collective threat. Conversely, groups related to gender identity or sexual orientation show a greater association with strategies of contempt, objectification, or moral degradation. Attacks directed at specific individuals, on the other hand, tend more often to manifest themselves in the form of generic insults or pathologization processes. This type of distribution suggests that offensive language is not random, but structured according to semantic and pragmatic patterns that reflect shared social representations.

The qualitative analysis of offensive spans also allowed us to identify several recurring patterns in the construction of the offense. Among these, dehumanization emerges as a particularly relevant strategy, through which the target is represented as inferior, animal, or devoid of human value. This is accompanied by pathologization, which consists in the attribution of characteristics of mental illness or abnormality, contributing to delegitimizing the subject. A further pattern is represented by the use of images of disgust, which evoke dirt, contamination or repulsion, reinforcing the emotional effect of the offense. Finally, forms of identity demonization are observed, in which social groups are associated with danger, violence, or collective threat. These strategies reflect dynamics well known in pragmalinguistic studies, according to which offensive meaning does not reside exclusively in literal content, but emerges from the interaction between context, inferences, and social stereotypes.

Exploratory and cross-sectional analyses also allowed us to evaluate the internal consistency of the annotated dataset. Overall, associations between *target_group* and *target_group_attribute* are semantically plausible, with a limited number of cases being ambiguous or difficult to classify. Such cases, often characterized by pragmatic ambiguity, ironic use, or implicit meanings, highlight the intrinsic complexity of the phenomenon and confirm the need for an interpretive, rather than purely lexical, approach in the classification of offensive language.

Overall, the results show how offensive language in the Italian dataset follows systematic and recognizable semantic patterns. Attack strategies vary de-

pending on the target, and identity categories play a central role in structuring the offense. In this context, adopting a semantic-pragmatic approach to annotation has proven particularly effective, allowing us to capture the regularities of the phenomenon more accurately than methods based exclusively on lexical correspondences.

5 Language Specific Analysis

The specific analysis for the Italian language allowed us to investigate deeper into the pragmatic and semantic characteristics of offensive language, highlighting phenomena that emerge only at a qualitative and contextual level.

One of the most relevant aspects concerns the strong dependence of the offensive meaning on the context of use. In many cases, in fact, offensiveness is not determined exclusively by the presence of taboo words, but by their discursive function within the utterance. As also highlighted in studies on taboo language, such expressions are often used with a connotative rather than referential function, helping to intensify the message or express evaluative attitudes (Graci, 2024). This phenomenon is particularly evident in the Italian dataset, where many offensive expressions take on different meanings depending on the pragmatic context.

A further distinguishing feature is the frequent use of indirect or implicit attack strategies. In several examples, the offense is not explicitly expressed through a direct insult, but emerges through inferences, allusions, or metaphorical constructions. This makes annotation particularly complex, as it requires an interpretation of the implicit meaning and shared knowledge between speaker and recipient. In such cases, approaches based on simple lexical correspondences are insufficient, while a semantic-pragmatic analysis allows us to grasp the real nature of the attack.

The qualitative analysis also highlighted the presence of ambiguity phenomena, often linked to ironic or sarcastic uses of language. Some expressions that, on a superficial level, could be interpreted as offensive, are actually toned down or reinterpreted in ironic contexts. Conversely, seemingly neutral statements can convey implicit offensive content. This type of ambiguity

represents a major source of difficulty for both automatic and manual annotation.

From an error analysis perspective, the main critical issues encountered concern precisely these borderline cases. In particular, errors or uncertainties in classification are observed when:

- The offense is implicit and does not contain obvious lexical indicators;
- The target of the attack is not clearly identifiable;
- The same expression can be associated with multiple semantic frames (e.g., overlap between disgust and dehumanization);
- The context is not sufficient to disambiguate the communicative intention.

A further difficulty concerns the distinction between attacks on groups and attacks on individuals. In some cases, expressions referring to a specific individual may implicitly recall broader categories, making the classification between individual and collective levels ambiguous.

Finally, it emerged that some semantic categories are intrinsically more difficult to distinguish than others. For example, the lines between dehumanization, disgust, and demonization can be blurred, as such strategies often share common semantic and rhetorical elements. This suggests that the categories used, while useful for analytical purposes, do not always represent rigidly separate classes, but rather interconnected dimensions of offensive language. Overall, the Italian specific analysis confirms that hate speech is a highly contextual and layered phenomenon, in which literal meaning, pragmatic implicatures, and sociocultural knowledge interact in complex ways. These findings reinforce the importance of interpretive and contextual approaches in analyzing the phenomenon, while highlighting the limitations of purely automatic methodologies or those based on predefined lexicons.

6 Cross Language Comparison Discussion

Although this study focuses exclusively on Italian language data, it is possible to place the results obtained within a broader comparative perspective,

referring to the existing literature on taboo and offensive language in multilingual contexts.

A first element of comparison concerns the universal nature of some properties of offensive language. Recent studies show that taboo expressions are generally characterized by high emotional activation and strong negative polarity, regardless of the language considered. The results obtained in this work are consistent with this perspective, highlighting how even in the Italian context, offensive language is based on recurring semantic strategies, such as dehumanization, disgust, and demonization.

However, alongside these similarities, the literature emphasizes the importance of sociocultural factors in determining the specific forms and categories of taboo language. In particular, what is considered offensive, as well as the perceived intensity of the offense, can vary significantly between different linguistic communities. In this sense, the patterns observed in the Italian dataset reflect not only general properties of offensive language, but also specific characteristics of the cultural context in question.

A relevant example concerns the distribution of semantic categories. In the Italian context, as highlighted by the analysis, categories related to social identity and group dynamics take on particular importance. This finding is consistent with studies showing that offensive expressions are often used to reinforce identity boundaries and inclusion/exclusion dynamics within a linguistic community. At the same time, the specific configuration of these categories may differ from that observed in other languages, depending on historical, social and cultural factors.

A further aspect concerns the differences in the semantic domains of taboo language. As highlighted in comparative studies, languages can privilege different domains (e.g., religious, sexual, or bodily) in the construction of offense. In the case of contemporary Italian, the analyzed dataset suggests a prevalence of strategies related to the body, sexuality, and social identity, rather than forms of religious blasphemy, which are less central than in other historical or linguistic contexts. This type of observation is consistent with linguistic analyses that highlight how the relative weight of different categories of taboo

language is subject to diachronic and cultural variations.

Finally, the results obtained confirm the difficulty of defining a universal taxonomy of offensive language. As highlighted in the literature, there is no univocal consensus on the number and nature of the categories that make up taboo language, precisely because of its strong dependence on the sociocultural context. Within this framework, the approach adopted in this study, based on flexible semantic pragmatic categories, proves particularly suited to capturing the complexity of the phenomenon, while maintaining a certain generalizability with respect to other linguistic contexts.

In summary, despite the absence of a direct multilingual empirical analysis, comparison with the literature suggests that offensive language simultaneously exhibits universal traits and cultural specificities. The findings regarding the Italian language fit into this framework, contributing to a broader and more nuanced understanding of the phenomenon.

7 Code Availability

The implementation is publicly available at: <https://github.com/FrancescaRomanaBlanda/Taboo-Language/tree/main>. There, the correspondent analysis for Sinhala and Korean can also be found.

8 Limitations

This study presents several limitations that should be considered when interpreting the results. First, the analysis is restricted to a single language (Italian), which limits the generalizability of the findings across different linguistic and cultural contexts. As offensive language is highly dependent on sociocultural norms, the observed patterns may not transfer directly to other languages.

Second, the annotation process relies partially on a language model-assisted approach. While this method allows for flexible and context-sensitive interpretation, it introduces potential biases and lacks a formal evaluation of annotation reliability. The absence of inter-annotator agreement metrics makes it difficult to quantify the consistency of the annotations.

Third, the dataset size is relatively limited and derived from pre-existing resources, which

may not fully capture the diversity of offensive language in natural settings. Additionally, the merging of datasets from different domains (tweets and news) may introduce heterogeneity that is not explicitly controlled for in the analysis.

Finally, the study primarily focuses on descriptive and exploratory analysis, without implementing predictive models or quantitative evaluation frameworks. As a result, the findings provide insight into the structure of offensive language but do not directly assess model performance or generalization capabilities.

9 Conclusion

This study analyzed offensive language in an Italian language dataset, with the aim of identifying its semantic patterns and underlying discursive strategies. By integrating existing resources and introducing semantic-pragmatic annotation, it was possible to enrich the dataset with information relating to target groups and attack categories, offering a more nuanced representation of the phenomenon.

The results show that offensive language is not a random set of expressions, but a structured system, in which attack strategies vary as a function of the target and the context. In particular, identity categories emerge as central elements in the construction of offense, while strategies such as dehumanization, pathologization, and demonization are recurrent and systematic. This confirms the importance of an approach that goes beyond the lexical level, privileging a semantic and pragmatic analysis of language.

From a methodological point of view, the use of language model assisted annotation has proven effective in capturing the complexity of the phenomenon, especially in the presence of implicit meanings, ambiguities, and contextual variations. Compared to alternative, more rigidly automated approaches, this method has allowed for greater interpretative flexibility and better annotation quality, making it particularly suitable for small datasets.

Despite the results obtained, the study has some limitations. First, the analysis focuses exclusively on the Italian language, limiting the

possibility of empirically generalizing the results to other linguistic contexts. Furthermore, the partly interpretative nature of the annotation introduces a margin of subjectivity, especially in ambiguous or ironic cases. Finally, the lack of more advanced quantitative analysis or systematic evaluation of annotation accuracy represents a further area that could be improved.

Future research could extend this work in several directions. In particular, it would be useful to apply the same annotation scheme to multilingual datasets, in order to systematically explore the differences and similarities between languages and cultures. Furthermore, the integration of more advanced computational models could allow part of the annotation process to be automated while maintaining sensitivity to semantic context. Finally, a more in depth analysis of pragmatic dynamics, for example, in relation to irony, sarcasm, or the creative use of offensive language, could contribute to an even more complete understanding of the phenomenon. In conclusion, this study helps highlight the complex, layered, and deeply contextual nature of offensive language, underscoring the need for flexible and multidimensional analytical approaches to its analysis.

References

- Roberto Graci. 2024. [Italian taboos and reporting practices: A sociolinguistic perspective](#). In Alessandro Capone, Roberto Graci, and Pietro Perconti, editors, *New Frontiers in Pragmalinguistic Studies: Theoretical, Social, and Cognitive Approaches*, volume 37 of *Perspectives in Pragmatics, Philosophy & Psychology*, pages 177–193. Springer Nature.
- Gesualdo La Porta. 2024. [Private information problem](#). In Alessandro Capone, Roberto Graci, and Pietro Perconti, editors, *New Frontiers in Pragmalinguistic Studies: Theoretical, Social, and Cognitive Approaches*, volume 37 of *Perspectives in Pragmatics, Philosophy & Psychology*, pages 229–243. Springer Nature.
- Brigid Maher. 2012. [Taboo or not taboo: swearing, satire, irony, and the grotesque in the english translation of Niccolò Ammaniti's Ti prendo e ti porto via](#). *The Italianist*, 32(3):367–384.
- Manuela Sanguinetti, Gloria Comandini, Elisa Di Nuovo, Simona Frenda, Marco Stranisci, Cristina Bosco, Tommaso Caselli, Viviana Patti, and Irene Russo. 2020. [HASPEEDE 2 @ EVALITA2020: Overview of the EVALITA 2020 Hate Speech Detection Task](#).
- Simone Sulpizio, Fritz Günther, Marco Marelli, and 1 others. 2024. [Taboo language across the globe: A multi-lab study](#). *Behavior Research Methods*, 56:3794–3813.